

C++ strings vs C-Strings

C-strings are **not first-class types** like the C++ **string** type. They do not work like the built-in types. Look at this example, which tries to **assign**, **compare** and **concatenate** two strings:

```
string str1 = "Hello", str2 = "World";
char cstr1[] = "Hello", cstr2[] = "World";

str1 = "Goodbye";           // assignment OK
cstr1 = "Goodbye";          // ILLEGAL
if (str1 < str2) ...         // comparison OK
if (cstr1 < cstr2) ...       // INCORRECT
str1 += ", ";               // OK
cstr1 += ", ";              // ILLEGAL
```

For the C++ **string** class, assignment, comparison and concatenation work in the same manner as the built-in types. Use the **assignment operator**, the **relational operators**, and the **+=**. **Not so** for C-strings, where you must use functions from the **<cstring>** header to perform the same functionality.

- **strcpy(dest, src)** is used instead of assignment
- **strcat(dest, src)** is used instead of **+=**
- **strcmp(cstr1, cstr2)** is used instead of the relational operators

In addition, in place of the member function **size()**, you use the **strlen(cstr)** function which counts the number of characters before the **'\0'**.



This course content is offered under a **CC Attribution Non-Commercial** license. Content in this course can be considered under this license unless otherwise noted.

C-String Assignment

Assignment means "copy the thing on the right into the storage on the left".

Instead of the assignment operator, used by the built-in types, C-strings use the `strcpy()` function, from the standard library header `<cstring>`, as shown below:



```
const size_t kMaxLen = 4096;
char dest [kMaxLen];
// Assume src is a C-style string
strcpy(dest, src);
```

Both `src` and `dest` are C-strings. (`src` is a common abbreviation for *source*, where the characters are copied *from*, while `dest` stands for *destination*, where the characters are copied *to*). `strcpy(dest, src)` copies the characters, one by one, from `src` into `dest`, stopping the `'\0'` is copied. However:

- You don't know if the **actual size** of the C-string source is less than **4095** characters (+1 for the null character). **Thus this code contains a security flaw.**
- You normally won't need anywhere near **4096** characters allocated for destination, so **the code is inefficient.**

It is up to you to ensure that there is enough space in `dest` to hold a copy of `src`. The icon used in front of the code does not mean that the code is buggy; instead, it means that the **function itself is intrinsically dangerous**; it's like the symbol found on rat poison.



The library function itself **makes no attempt to check whether the destination has enough room** to hold a copy of the source string. Even if there is not enough memory the function keeps copying, possibly overwriting other data; this called a **buffer overflow**.



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

The strncpy Function

The possibly safer `strncpy()` function copies only a specified number of characters from `src` to `dest`. Here's how it works:

- Call `strncpy()` with a `dest`, a `src`, and a `count` of characters.
- If the `'\0'` in `src` is found **before** the specified number of characters have been copied, then `strncpy()` will fill the remainder with `'\0'`.
- If the **NUL** character is not found in `src` before the number of characters have been copied, then **you must manually append a terminating NUL**.

Here is a **semi-safe copy**, given the previous example that avoids overflow (although it doesn't ensure that all of `src` was actually copied; for that you need a loop).

```
dest[kMaxLen - 1] = '\0';    // pre-terminate
strncpy(dest, src, kMaxLen - 1);
```

If I've seemed somewhat equivocal about using `strncpy`, you should know that it's widely regarded as a still unsafe function. If you want to know more, here are some links:

- [Stop Using strncpy already!](#)
- [strncpy? Just say no](#)
- [strncpy: Not the function you are looking for](#)



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

The strcat Functions

Concatenation is the province of the `strcat()` (completely unsafe), and the `strncat()` (marginally safer) functions. Here is a (buggy) example using the functions:



```
const size_t kLen = 10;
char cstr[kLen] = "Goodbye";
strcat(cstr, " cruel world!"); // OOPS
cout << strlen(cstr) << " " << cstr << endl;
```

When you run, you'll likely see:

```
Goodbye cruel world!
```

The C-string `cstr` has room for 9 characters, but you **appear to** have stuffed **21** characters (including the `NUL`), into that smaller space. Not really, of course: **this is a buffer overflow** and the actual **results are undefined**.

The `strncat()` function is marginally safer, if **fairly tricky to use correctly**. If used incorrectly, it overflows just like `strcat()`. Here is the prototype:

```
char * strncat(char *dest, const char *src, size_t count);
```

The tricky part is that `count` is not the maximum size of the result, but the maximum number of characters to be copied; you must first calculate the **correct combined maximum**, before calling the function.

```
const size_t kLen = 39; // max total characters
const cstr[kLen + 1] = "This is the intial string";
const char *str2 = "Extra text to add to the string";
strncat(cstr, str2, kLen - strlen(cstr));
```

This **isn't efficient** (since you need to count the characters in `cstr` first), but it **does stop copying when the destination string is full**.

Security Note: `strncat()` does not check for sufficient space in `dest`; it is therefore a potential cause of buffer overruns. Keep in mind that `count` limits the number of characters appended; it is not a limit on the size of `dest`.



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

Comparing C-Strings

Do not use the relational operators (<, ==, etc.) to compare C-strings. Instead, use the library function `strcmp()`, which compares `s1` and `s2` **lexicographically** and returns an integer indicating their relationship:

- **Zero** if the two strings are equal.
- **Negative** if the first string lexicographically precedes the second string. (Lexicographically simply means "in dictionary order").
- **Positive** if the first string lexicographically follows the second string.

To use `strcmp()` correctly:

- Call the function and save the int it returns.
- Use the returned value with a relational operator.
- **Don't** treat the return value from `strcmp()` as a Boolean expression.
- Don't repeatedly call `strcmp()` on the same strings (inefficient).

Here's a quick example. The C-strings `s1` and `s2` are initialized elsewhere. Since we don't need to modify either argument, we can use "pointer-style" C-strings.

```
const char *s1 = ..., *s2 = ...;

int result = strcmp(s1, s2);

if (result == 0) ...           // equal
else if (result < 0) ...       // s1 < s2
else ...                       // s1 > s2
```



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

The strlen Function

In this lesson we're going to look at several implementations of the standard library functions beginning with `strlen()`. We'll finish by learning to write your own C-String functions.

To find the length of a string, you **count characters** until you reach the `'\0'`. Here is an implementation that uses array notation.

```
size_t strlen(const char str[])
{
    size_t len = 0;
    while (str[len] != '\0') len++;
    return len;
}
```

*Note that the return type **must be size_t** (not **int**), because we can't have a negative length on a string. The array must be **const**, otherwise it **would be illegal** to call the function using a C-string literal.*

Another alternative is to advance the pointer until it reaches the end of the string, and then to **use pointer subtraction** (or **pointer difference**) to determine the number of characters. Here's a version that does that:

```
size_t strlen(const char *str)
{
    const char *cp = str;
    while (*cp != '\0') cp++;
    return cp - str;
}
```

We can actually write this in an even more cryptic style. I **don't encourage you to write code like this**, since it is quite a bit more error prone, but it **is a common C idiom** so you should recognize it when you see it.

```
size_t strlen(const char *str)
{
    const char *cp = str;
    while (*cp++) /* do nothing */;
    return cp - str - 1; // cp points to 1 past the NUL
}
```



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

The strcpy Function

The `strcpy()` function is often **even more cryptic** than `strlen()`.

```
char * strcpy(char *dest, const char *src)
{
    char *result = dest;
    while (*dest++ = *src++) /* do nothing */;
    return result;
}
```

This **very, very common idiom** has so many potential pitfalls, that it is likely that your IDE will mark it with a warning. Although technically not incorrect, it is intrinsically dangerous code, since a small mistake can break the loop entirely.

- The **body of the while loop is empty**; all of the work occurs in the extremely streamlined test expression: `*dest++ = *src++`
- This expression is **not a comparison**, but an **embedded assignment**. If you accidentally use a comparison, the loop will not work.
- The expression copies the character addressed by `src` into the address indicated by `dest`, incrementing each pointer after the character is copied. **If you use prefix increment instead of postfix, this does not work**.
- The result is zero—and therefore **false**—only when the code copies the **NUL** character at the end of the string.

Note that this leaves both pointers pointing one-past the **NUL** characters in their respective strings.



This course content is offered under a **CC Attribution Non-Commercial** license. Content in this course can be considered under this license unless otherwise noted.

The strcmp Function

Like `strcpy()`, most implementations of `strcmp()` are cryptic. Here's the version from GNU C:

```
int strcmp(const char *s1, const char *s2)
{
    const unsigned char *a1, *a2;
    for (a1 = reinterpret_cast<const unsigned char *>(s1),
         a2 = reinterpret_cast<const unsigned char *>(s2);
         *a1 == *a2; a1++, a2++)
        if (*a1 == '\0') return 0;
    return *a1 - *a2;
}
```

The GNU version of `strcmp()` returns **the difference** between the first two mismatched characters. `a1` and `a2` are temporary pointers to **unsigned char**, so the characters can be interpreted as raw values between **0-255**. The pointers are initialized by using a **reinterpret_cast**.

Here is an alternate (Apple/Next/PPC) version of the same function, which returns **0, +1** and **-1** instead of the difference between the characters. This version, written in 1992, uses traditional C-style casts to handle the **signed/unsigned** instead of a C++ **reinterpret_cast**.

```
int strcmp(const char *s1, const char *s2)
{
    for ( ; *s1 == *s2; s1++, s2++)
        if (*s1 == '\0') return 0;    // reached both NULs. Equal
    return ((*((unsigned char *)s1)
            < *((unsigned char *)s2) ? -1 : +1);
}
```



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

Writing Your Own Functions

To write your own C-String functions you can use either array notation or pointer notation, whichever you find more convenient; **neither** is more efficient than the other. The things you need to remember are:

- **Find** the **NUL** character in the string. All C-String functions rely on this.
- **Preserve** the **NUL** character in the string. It is up to you to make sure that any destination strings are correctly terminated.

To make this more concrete, let's look at a couple of examples.



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

Find First

To find the first occurrence of a particular character in a string, you'd employ the **linear search** algorithm:

```
Loop through a string until the NUL character
If current character is the target
    Return its index
Return the error code
```

Assuming that we use **-1** for the error code then an **array-notation** implementation of the function could look like this:

```
int find(const char a[], char target)
{
    for (int i = 0; a[i] != '\0'; ++i)
        if (a[i] == target)
            return i;
    return -1;
}
```

A (more cryptic) **pointer-notation** implementation might look like this:

```
int find(const char* s, char target)
{
    auto *p = s;
    while (*p && *p != target) p++;
    if (*p) return p - s;
    return -1;
}
```

The **temporary pointer** **p** is moved through the C-string **s**. The expression ***p** is false when the **NUL** is encountered. Since the loop **must end** when you encounter the **NUL**, or, when you find the **target**, you know that the loop **terminates in every case**.

After the loop is over, there are **two** possibilities. If **p** is pointing at **any** character, it **must** be the **target** character. That means you can use **pointer difference** to return the index. Otherwise, **p must** be pointing at the **NUL** character and you can return **-1**.



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

Find Last

You might think that the easiest thing would be to start at the back of the **string** and then loop towards the front. That's what you'd do with a C++ **string**. However, with C-strings, you can't find the length **without first looking at every character**, so looping backwards is actually more inefficient than simply going forward, saving the position each time the target is found.

Here's an efficient **array-notation** implementation of the function:

```
int find_last(const char a[], char target)
{
    int result = -1;
    for (int i = 0; a[i] != '\0'; ++i)
        if (a[i] == target)
            result = i;
    return result;
}
```



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

Find First of Any

Suppose you want to find the position of the first digit inside a C-string. You can't just use `find()` since you want to look for **any** digit. You'd want a function you could call like this:

```
int pos = first_of_any(cstr, "0123456789");
```

Here's the algorithm you use:

*Look through every character in str
Compare the character to every character in target
If found return the index (in str)
return error code*

Here's an implementation of this algorithm:

```
int first_of_any(const char *str, const char* target)
{
    const char * p1 = str;
    while (*p1 != '\0')
    {
        const char * p2 = target;
        while (*p2 != '\0')
        {
            if (*p2 == *p1) return p1 - str; // found it
            p2++;
        }
        p1++;
    }
    return -1;
}
```



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

Finding Substrings

Searching for a C-style substring inside another C-string is a little bit of work.

Similar to `first_of_any()` this is most easily done by using three temporary pointers.

Here's the algorithm:

```
String str, string target
Pointer p set to str
While *p != 0
    Pointer p1<-p
    Pointer p2<-target
    While *p1 && *p2 && *p1 == *p2
        p1++, p2++
    If *p2 == '\0' return p - str
    p++
Return the error code
```

You can implement it yourself on the next page.



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

CodeCheck

Reset