


# The Wall of Abstraction

A class is an interface paired with an implementation, similar to the `Time` structure you created in the lesson on Information Hiding. This is called the **wall of abstraction**, illustrated by the comic below.



The **public interface** specifies how clients interact with objects, and the **private implementation** specifies how the functions in the public interface are implemented.

The `Time` struct (even when paired with an interface, so the data is hidden), still allows users to **directly manipulate** its data members. Classes take a different approach; with classes, the data inside your objects will **only be accessible by the member functions**, **forcing** the client to access and modify data in a safe way.



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# Class Definition Syntax

---

## ► The Time Class Definition

To create a **class definition** for **Time**, similar to the structure from last week, follow these rules:

1. Instead of **struct**, use the **class** keyword. There is no **public** in front of this as with Java.
2. The **public** keyword, followed by a colon, indicates the start of the **public interface**. Here we **prototype** the member functions it contains.
3. The member functions **hours()**, **minutes()**, **sum()**, **difference()** and **write()**, all access the **hours** and **minutes** data members **without changing them**. When this is the case, add the **const** keyword **after the argument list**. We say these functions are **accessors**.
4. The **read()** member function does **modify** the **Time** object. This is called a **mutator**.
5. The class definition **ends with a semicolon**, just like a structure. This is not optional.

## Data Members

Most of the implementation will appear inside a **.cpp** file. **Defining the data members** which store **object state**, is **written inside the header file instead**. A common practice is to use a special indicator like **m\_** to show that it is a data member.

The **Time struct** used two individual data members: one for **hours** and one for **minutes**. This is fine; it allows you to store all of information needed. By adding **private**, you can **prevent clients of Time from accessing the data members directly**.

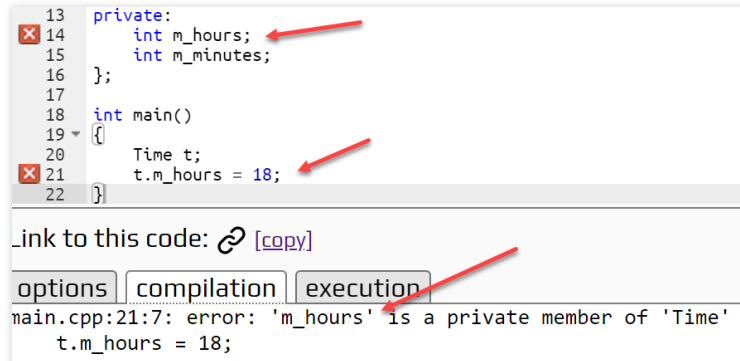


This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# Public and Private

So, what do **public** and **private** mean in C++? If a member of a class is **public**, then **any part of your code** can access and manipulate it directly. If you have a **public** member function, any code can **call it using an object of that type**. If a data member is marked **private**, then **only** member functions of the class can access it.

The **public** and **private** keywords are the C++ mechanism for **defining interfaces and enforcing encapsulation**. Once you add **private**, the compiler enforces the **appropriate encapsulation**.



```
13 private:
14     int m_hours;
15     int m_minutes;
16 };
17
18 int main()
19 {
20     Time t;
21     t.m_hours = 18;
22 }
```

Link to this code: [\[copy\]](#)

options compilation execution

main.cpp:21:7: error: 'm\_hours' is a private member of 'Time'  
t.m\_hours = 18;

By prohibiting clients from directly accessing **private** data, the implementation can assume that all access to that data goes through the **public** interface (unlike the **Time struct** of last week, where clients **should use the member functions**, but **were not prohibited** from directly accessing the data members **m\_hours** and **m\_minutes**.)

*Actually, the only **real** difference between **class** and **struct** in C++ is that with a **struct**, the members are **public** by default; with a **class** they are **private**. By convention, we will use **struct** for **POD** (plain-old-data) data types, and **class** for encapsulated types.*



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# The Implicit Parameter

Consider the implementation of the `hours()` member function of the `Time` class shown here:

```
int Time::hours() const
{
    return m_hours;
}

int main()
{
    Time t;                // a Time object
    cout << t.hours() << endl; // value of t::m_hours
}
```

The `hours()` member function **does not** contain a local variable named `m_hours`. But, the function still compiles and runs correctly. Why?

In a **member function**, you may **directly access and manipulate** any or all of the class's **data members** by referring to them by name. You don't need to indicate that `m_hours` is a data member, nor do you specify **which** `Time` object it belongs to.

C++ assumes that all data members **are the data members of the receiver object**, and so the line `return m_hours` means “return the value of the `m_hours` data member of the object on which this function was invoked.” In such a case, **the receiver object is known as the implicit parameter**, passed to every member function.



This course content is offered under a [CC Attribution Non-Commercial](#) license. Content in this course can be considered under this license unless otherwise noted.

# The Pointer this

Behind the scenes, the implicit parameter is a **pointer to the calling object**. Every member function has an implicit parameter. Thus the effective signature for the `hours()` function is as if you had declared it like this:

```
int hours(const Time* const this);
```

The keyword `this` is the **name which is automatically supplied** for the implicit parameter. The `const` following the `Time*` means that the value inside the pointer can never be changed; it always points to the block of data containing the object's data members. The `const` following the member function header means that the implicit parameter is a pointer to a `const Time` object.

If you wish, you can **explicitly** use the pointer when calling other member functions, or accessing data members:

```
int Time::hours() const
{
    return this->m_hours;
};
```

## Initializing this

When you **call a member function** like this:

```
Time t;           // a Time object
cout << t.hours() << endl;           // value of t::m_hours
```

That call is **implicitly translated** into code that acts as if you had written:

```
Time t;           // a Time object
cout << hours(&t) << endl;           // value of t::m_hours
```

Because of this call, the `this` pointer is initialized to the **address of the calling object**.



This course content is offered under a [CC Attribution Non-Commercial](https://creativecommons.org/licenses/by-nc/4.0/) license. Content in this course can be considered under this license unless otherwise noted.

# The sum Member Function

Let's examine the behavior of **this** and **const** a little more closely, by considering the **sum()** member function from **Time**:

```
class Time
{
public:
    Time sum(const Time& rhs) const;
    . . .
};
```

When you add two **Time** objects (**a + b**) together like this:

```
Time after = a.sum(b);
```

The **caller** (the implicit parameter) is the left-hand-side of the expression **a + b**. Thus, the effective implicit prototype for the function is similar to this:

```
Time sum(const Time* lhs, const Time& rhs);
```

In the implementation, however, instead of the **explicit lhs parameter** shown here, you'd use the keyword **this** to access the data members.

```
Time Time::sum(const Time& rhs) const
{
    auto tMinutes = this->m_hours * 60 + this->m_minutes;
    auto dMinutes = rhs.m_hours * 60 + rhs.m_minutes;
    . . .
}
```

If you leave off the keyword **this**, it is assumed. Notice that when you implement a **const** member function, you **repeat** the word **const** in the implementation.



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# Setters

---

In Java, many classes have member functions that start with `set`. These are called **mutators**, since they **change the state** of the object. Mutators should **validate data** written to the object to **enforce the class invariants**.

With properly written mutators, the errors described in earlier lessons **cannot occur**. Consider your `Time` class. If you were to add `setHours()` and `setMinutes()` members to the class, you would have to enforce these restrictions:

- `m_hours` must be between `0` and `23` inclusive.
- `m_minutes` must be between `0` and `59` inclusive.

Unlike the `read()` member function, where you could put the stream into a failed state, if these conditions were not met, in a mutator you need to **throw** an exception like this:

```
void Time::setHours(int h)
{
    if (h < 0 || h > 23) throw out_of_range("...");
    m_hours = h;
}
```



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# Getter & Setter Patterns

---

The pattern of pairing a "getter" along with a "setter" function is common, and you will see it in any major C++ project you work on. Unlike Java, the actual `get*` and `set*` name pattern is not as common. Instead, what programmers often do is write a pair of overloaded functions.

Instead of the name `getHours()` or `setHours()`, use the name `hours()` for both of them.

- The accessor is `const` and returns a value.
- The non-`const` mutator returns a reference, which can be assigned to.

```
Time Time::hours() const { return m_hours; } // accessor
Time& Time::hours()      // mutator
{
    if (h < 0 || h > 23) throw out_of_range("...");
    return m_hours;
}
```

In general, it is safer to allow clients to **read the values** of the data members than it is allow them to **change** those values. As a result, "setters" are far less common than "getters" in class design. Classes with no mutators at all, are called **immutable classes**.



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.



# Constructors

Initializing object data is the responsibility of the **constructor**, which always has **the same name as the class** and which **never has a return type**.

```
class Time
{
public:
    Time();    // default constructor
    ...
};
```

A constructor is a member function which **initializes an object into a well-formed state** before clients start manipulating it. When C++ creates an object from a class:

1. It **allocates a block of memory** large enough to store the data elements.
2. It passes **the address of that block** of memory to the constructor function. The address is the `this` pointer inside the constructor function.

The constructor **is called automatically** whenever an object is created. If you have a class that defines a constructor, that constructor is **guaranteed to execute** whenever you create an object of the class type.

## Default Constructors

The **default constructor** is the constructor which takes **no arguments** and which should initialize **all of its data members** to an appropriate **default** value. Alternatively, since C++11, you may **provide an initial value** when defining the data members, just as in Java.

If you do not provide a constructor, the compiler will "write" one for you. This is called the **synthesized default constructor**. If you use in-definition initializers, then this is perfect.



This course content is offered under a [CC Attribution Non-Commercial](https://creativecommons.org/licenses/by-nc/4.0/) license. Content in this course can be considered under this license unless otherwise noted.

# Member Initialization

In C++, all constructors must initialize **all primitive types**. A C++ constructor does not need to initialize any object members (like **string** or **vector**).

This is **exactly the opposite from Java**, where you must initialize all of the object instance variables, or they are set to **null** (an invalid object). In Java, all primitive instance variables are automatically initialized to **0** like this:

```
public class Point
{
    private String name;
    int x, y;
    public Point() {}
}

Point p = new Point(); // x,y->0, name is null (invalid)
```

In C++, if you fail to initialize a primitive data member, then it assumes **whatever random value** was in memory; if you don't initialize an object, such as **string** or **vector**, its default constructor will **automatically** run, and it is still a valid object.

```
class Point
{
public:
    Point() {}
private:
    string name;
    int x, y;
};

Point p; // x,y->random, name is valid empty string
```

Of course, if you provide in-definition initializers for your primitive data members, they **will** automatically be initialized, even if your construct does not explicitly initialize them.



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# Working Constructor

With the `Time` class we might like to have another, overloaded constructor which takes hours and minutes. This is generally known as the **working constructor**. Here is the public interface of `Time` with both of these constructors.

```
class Time
{
public:
    Time();           // default
    Time(int h, int m); // working
    . . .
};
```

Unfortunately, if you have **any** explicit constructors, the synthesized one **is deleted**, so you have to add an **explicit default constructor** as done here. In C++11, however, you can just add the phrase `=default;` to the end of the prototype in the class header, and the compiler will **retain** the synthesized constructor that it normally writes.

The **implementation** of the constructors goes into the `.cpp` file along with the other member functions. The job of the constructor is to **initialize the data members**, so in the `Time` class, you might have code that looks something like this.

```
Time::Time() { m_hours = m_minutes = 0; }
Time::Time(int hours, int minutes)
{
    m_hours = h;
    m_minutes = m;
}
```



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# Assignment vs. Initialization

Before we talk about constructors, look at these two statements:

```
1 | string a = "Bob", b;      // initialization
2 | b = "Bill";              // assignment
```

1. Two **string** objects are created and initialized on line one; **a** is initialized using the C-String **"Bob"**, and **b** is initialized to the empty **string** by running the default constructor.
2. The **string** object **b** is destroyed (its destructor is run), a new **string** object is initialized with **"Bill"**, and that new **string** object replaces the **string** object originally held by **b**.

The variable **b** is first initialized, then destroyed, then assigned. **This is inefficient.**

## Assignment in a Constructor

The body of the constructor is executed **after** the data members have been initialized. You may use **assignment** to place a new value into these data members. For primitive types, the cost of doing this is negligible, but for object types, such assignments mean that **data members are constructed twice**—once at initialization and once at assignment. Here's an example. (The implementation is inline to shorten the code.)

```
class Person
{
public:
    Person(const string& name) { m_name = name; }
private:
    string m_name;
};
```

When you write **Person p("Fred")**, the **m\_name** data member first calls the default constructor to create an empty **string** object. Then, in the body of the constructor, the default-constructed **string** is destroyed when assigning **name** to **m\_name**. **This is inefficient**, and you want to avoid it.



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# The Initializer List

---

We can instruct the compiler to initialize the individual data members before the body of the constructor is entered instead. This is called the **initializer list**:

- It follows the parameter list and is preceded by a colon ( : )
- It is followed by a list of member names and their initializers.
- Initialization occurs **in the order the members are declared in the class**.

In C++98 the initializers are placed in parentheses; in C++11 use either parentheses or braces. You cannot use the assignment operator. Here is the same class using the initializer list. In this case, the **name** data member is **only constructed once**:

```
class Person
{
public:
    Person(const string& name) : m_name(name) { } // empty body
private:
    string m_name;
};
```



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# The Working Constructor

---

As you saw in the last lesson, **working constructor** is the short-hand description of a constructor that takes as many user-supplied arguments as possible. In the `Time` class the working constructor looks like this:

```
Time(int hours, int minutes);
```

In the `.cpp` file, you might have code that looks something like this, using the initializer list, also from the last lesson:

```
Time::Time(int hours, int minutes)
: m_hours(hours), m_minutes(minutes)
{
    // validate the constructor arguments
    assert(hours >= 0 && hours < 24);
    assert(minutes >= 0 && minutes < 60);
}
```

It is important that your constructor initializes your object so that it is **in a valid state**. In the example shown here, we've used `assert()` on the assumption that it is a programming error if an invalid `Time` is constructed. If, however, you were constructing `Time` objects using external data, it is possible you would want to **throw** and **catch** an exception instead.



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.

# Conversion Constructors

---

A **conversion constructor** is a constructor (usually 1-argument) that **implicitly converts between one type and another**. Here's an overloaded conversion constructor that converts between fractional hours and hours and minutes:

```
Time(double hours);
```

The implementation of the constructor (converting first to seconds, then extracting the hours and minutes) could look like this:

```
Time::Time(double hours)
{
    assert(hours >= 0 && hours < 24);
    const int kSecondsPerHour = 3600;
    const int kHoursPerMinute = 60;
    auto time = static_cast<int>(hours * kSecondsPerHour);
    m_hours = time / kSecondsPerHour;
    m_minutes = time % kSecondsPerHour / kHoursPerMinute;
}
```



This course content is offered under a [CC Attribution Non-Commercial](https://creativecommons.org/licenses/by-nc/4.0/) license. Content in this course can be considered under this license unless otherwise noted.

# Implicit vs. Explicit

---

Conversion constructors can be **implicit** (which is the default), or **explicit**.

The implicit conversion constructor is called any time the compiler needs a **Time** object, but finds a **double** that it can convert. Consider this fragment of code:

```
1 | Time bed_time(23, 30);    // 11:30
2 | bed_time = 5.2;          // WHAAAAAT?
```

You would **expect** that line 2 would be a syntax error, but, surprisingly, it is not. Instead, **the conversion constructor is silently (implicitly) called**, and **bed\_time** is changed to **5:12 am**. Probably not what you expected.

You can add **explicit** as a modifier to the prototype to prevent this:

```
explicit Time(double hours); // 7.51 -> 7:35
```

The keyword **explicit** **only** goes in the class definition. It **is not repeated** in the **.cpp** file. Sometimes, as you'll see when you cover symmetric overloaded operators in CS 250, you'll **want** to allow implicit conversion. For instance the **string(const char\*)** constructor is **not explicit**. Most of the time, however, **explicit** is preferred.



This course content is offered under a CC Attribution Non-Commercial license. Content in this course can be considered under this license unless otherwise noted.